

GRLite Sandbox v39: per-particle self-field subtraction with empirical $\epsilon(v_x, v_y)$ correction

Daniel Marks

2026-06-01

This document is an addendum to `gr_sandbox_v38.tex`. All prior content remains current; the v39 work introduces an OPTIONAL per-particle field subtraction mechanism for clean self-force cancellation, with a velocity-dependent residual-correction knob $\epsilon(v_x, v_y)$ intended for analytic or empirical calibration.

1 Motivation: HE adjoint fails for moving sources under FDTD

The Hockney–Eastwood adjoint condition guarantees zero spurious self-force on a STATIC source when the deposit kernel D and gather kernel G satisfy $G = D^T$, provided the field is solved quasi-statically (Poisson, $c \rightarrow \infty$). In that limit the field is an instantaneous functional of the current source positions and the symmetry argument that gives $F_{\text{self}} = 0$ closes.

For a particle moving at velocity $\vec{v}$ in a wave-evolved FDTD scheme the condition no longer applies. The discrete leapfrog

$$\phi_{i,j}^{n+1} = 2\phi_{i,j}^n - \phi_{i,j}^{n-1} + (c \Delta t)^2 (\nabla_d^2 \phi_{i,j}^n + s_\phi \rho_{i,j}^n) \quad (1)$$

is NOT Lorentz invariant: the field at any time is built from the source at retarded times along characteristic cones of the discrete operator, and the steady-state response to a uniformly-moving source is NOT the boosted continuum Coulomb shape. The residual gradient at the particle’s current position is the discrete spurious self-force $F_{\text{spurious}}(\vec{v})$.

Empirical findings (Stage 37e–37g, all equal-distance methodology with mass = 1, BUMP $R = 8$):

- $F_{\text{spurious}}(\vec{v})$ is non-monotonic with sign changes at $v \approx 0.045$ and $v \approx 0.075$.
- The threshold velocities are INDEPENDENT of the deposit kernel (TSC bare, TSC + $N = 4/16/64$, BUMP $R = 2.5/4/6/8$ all give the same sign at every v in the scan). Kernel choice changes only magnitude.
- The threshold pattern lives in the static $-q\nabla\phi$ channel; turning off $-q\partial_t A$ and $+q\vec{v} \times \vec{B}$ does not move the thresholds.
- EM (repulsive self, $c_{\text{em}} = +4\pi k_e$) and GR (attractive self, $c_{\text{grav}} = -4\pi G_{\text{eff}}$) give F_{spurious} of identical magnitude and OPPOSITE sign at every v .

These observations are consistent with the HE adjoint being a quasi-static theorem: kernel matching cannot save you once the field carries information from retarded source positions.

2 Proposed mechanism: optional per-particle self-field

2.1 Architecture

Each particle continues to deposit into the collective field as today. In addition, an OPTIONAL per-particle field set may be allocated for particle p :

$$\{\phi_g^{(p)}, A_{gx}^{(p)}, A_{gy}^{(p)}, \phi_{em}^{(p)}, A_x^{(p)}, A_y^{(p)}\} \quad (2)$$

sourced ONLY by particle p 's deposit. The collective and per-particle fields evolve independently under the same leapfrog (equation 1); both see the same damping/PML treatment so their wakes match.

At gather time, the force on particle p is

$$\vec{F}_p = \text{gather}(\phi^{\text{collective}}, \vec{x}_p) - (1 + \vec{\epsilon}(\vec{v}_p)) \cdot \text{gather}(\phi^{(p)}, \vec{x}_p) \quad (3)$$

where $\vec{\epsilon}(\vec{v}_p) = (\epsilon_x(\vec{v}_p), \epsilon_y(\vec{v}_p))$ is a per-axis modulation (the v -dependence may be anisotropic; see section 3).

Interpretation of $\vec{\epsilon}$:

- $\vec{\epsilon} = \vec{0}$: full subtraction. $\vec{F}_p$ contains only contributions from OTHER particles' fields. Self-force is identically zero (modulo float roundoff and warmup transients).
- $\vec{\epsilon} = -\vec{1}$: no subtraction, $\vec{F}_p$ reduces to the current PIC behavior (collective field gather only).
- $\vec{\epsilon}$ between 0 and -1 : partial self-coupling.

By construction, when the per-particle field is OFF for particle p , the gather collapses to the current behavior ($\vec{F}_p = \text{gather}(\phi^{\text{collective}})$), so existing tests are unaffected unless they explicitly opt in.

2.2 API sketch

```
gr_sim_particle_enable_self_field(sim, idx);
gr_sim_particle_disable_self_field(sim, idx);
gr_sim_particle_set_self_field_epsilon(sim, idx, eps_x, eps_y);
gr_sim_particle_set_self_field_epsilon_fn(sim, idx, callback);
int gr_sim_particle_has_self_field(sim, idx);
```

2.3 Cost

For N particles with self-field enabled and a $W \times H$ grid:

- Memory: $6 \cdot 3 \cdot W \cdot H \cdot \text{sizeof(float)}$ per opted-in particle. At $W = H = 256$ that is ~ 4.7 MB per particle. A 2-particle binary uses ~ 9 MB extra.
- Compute: each FDTD step performs $1 + N_{\text{opt}}$ leapfrog updates per field component. Linear in the number of opted-in particles.
- Esirkepov deposits run twice per opted-in particle (once into collective, once into private). Same kernel, same continuity properties.

For GRLite's pedagogical scope ($N \lesssim 10$), this is well within budget.

3 Toward an analytic $\vec{\epsilon}(\vec{v})$

The goal of $\vec{\epsilon}(\vec{v})$ is to make the residual self-force vanish at constant $\vec{v}$ even when the per-particle self-field deposit/evolution doesn't exactly mirror the collective contribution (warmup transients, anisotropy in the gather stencil, etc.). Ideally $\vec{\epsilon}$ would be derivable analytically from the discrete dispersion relation.

3.1 Dispersion relation for the leapfrog wave equation

Plug a spatially periodic plane-wave ansatz into equation 1. With $\phi_{i,j}^n = \hat{\Phi} \exp(i(k_x i \Delta x + k_y j \Delta x - \omega n \Delta t))$ the time-stencil gives

$$-4 \sin^2(\omega \Delta t / 2) \hat{\Phi} = (c \Delta t)^2 \left[\tilde{\nabla}_d^2(\vec{k}) \hat{\Phi} + s_\phi \hat{\rho} \right] \quad (4)$$

where the 5-point Yee Laplacian's symbol is

$$\tilde{\nabla}_d^2(\vec{k}) = -\frac{4}{\Delta x^2} [\sin^2(k_x \Delta x / 2) + \sin^2(k_y \Delta x / 2)] \equiv -\frac{4}{\Delta x^2} S^2(\vec{k}). \quad (5)$$

The discrete dispersion relation for the source-free wave equation is therefore

$$\sin^2\left(\frac{\omega \Delta t}{2}\right) = \left(\frac{c \Delta t}{\Delta x}\right)^2 S^2(\vec{k}) \equiv r^2 S^2(\vec{k}), \quad (6)$$

where $r = c \Delta t / \Delta x$ is the CFL number. In the continuum limit $\Delta t, \Delta x \rightarrow 0$ one recovers $\omega = c|\vec{k}|$. The discrete RHS is anisotropic: at fixed $|\vec{k}|$, ω is maximal along $\hat{x}$ or $\hat{y}$ and minimal along the diagonal.

3.2 Steady-state response to a uniformly-moving source

A uniformly-moving deposit translates rigidly with $\vec{v}$, so its Fourier representation is supported on the worldline $\omega = \vec{k} \cdot \vec{v}$:

$$\hat{\rho}(\vec{k}, \omega) \propto \hat{S}(\vec{k}) 2\pi \delta(\omega - \vec{k} \cdot \vec{v}), \quad (7)$$

where $\hat{S}(\vec{k})$ is the deposit kernel's spectrum. Solving for $\hat{\Phi}$ at the source's worldline frequency:

$$\hat{\Phi}(\vec{k}) = \frac{(c \Delta t)^2 s_\phi q \hat{S}(\vec{k})}{4[r^2 S^2(\vec{k}) - \sin^2(\vec{k} \cdot \vec{v} \Delta t / 2)]}. \quad (8)$$

The denominator is the discrete analog of $c^2 k^2 - (\vec{k} \cdot \vec{v})^2$. In the continuum limit it factorizes and the integral that follows vanishes by symmetry; on the grid neither holds.

3.3 Self-force as a Brillouin-zone integral

The gather operation, assumed HE-adjoint ($G = D$, same kernel spectrum $\hat{S}$), produces a force on the particle equal to

$$\vec{F}^{(\text{self})}(\vec{v}) = -q \int_{\text{BZ}} \frac{d^2 k}{(2\pi)^2} i \vec{k} |\hat{S}(\vec{k})|^2 \hat{\Phi}(\vec{k}) \quad (9)$$

with $\hat{\Phi}$ from equation 8 and the integral over the first Brillouin zone $k_\alpha \in (-\pi/\Delta x, \pi/\Delta x]$.

In the continuum, the integral is over all of $\mathbb{R}^2$ and the integrand is odd in $\vec{k}$ (after the boost) so it vanishes by symmetry: $\vec{F}_{\text{cont}}^{(\text{self})} = 0$. On the grid, the BZ is compact and the lattice's $\hat{x}/\hat{y}$ vs. diagonal anisotropy breaks the rotation symmetry of the continuum boost; the integral is generically nonzero and is a function of $\vec{v}$:

$$\boxed{\vec{F}_{\text{discrete}}^{(\text{self})}(\vec{v}) = -\frac{q^2 s_\phi (c \Delta t)^2}{4} \int_{\text{BZ}} \frac{d^2 k}{(2\pi)^2} \frac{i \vec{k} |\hat{S}(\vec{k})|^2}{r^2 S^2(\vec{k}) - \sin^2\left(\frac{(\vec{k} \cdot \vec{v}) \Delta t}{2}\right)}} \quad (10)$$

3.4 Why the integral fails to vanish on the grid

In the continuum, the denominator factors as $(c|\vec{k}| - \vec{k} \cdot \vec{v})(c|\vec{k}| + \vec{k} \cdot \vec{v})$. The integrand is then odd under $\vec{k} \rightarrow -\vec{k}$ after a suitable rotation aligning $\hat{z}$ with $\vec{v}$, and the integral vanishes. On the discrete grid, two distinct mechanisms break that cancellation:

Anisotropy of $S^2(\vec{k})$. The lattice Laplacian's symbol $S^2(\vec{k}) = \sin^2(k_x \Delta x/2) + \sin^2(k_y \Delta x/2)$ is NOT rotation-invariant. At fixed $|\vec{k}|$ it is maximal along $\hat{x}$ or $\hat{y}$ (where one term saturates near $\sin^2(\pi/2) = 1$) and minimal along the diagonal. A continuum Lorentz boost rotates the integration measure isotropically; the lattice does not respect that rotation, so a boost mixes $\hat{x}/\hat{y}$ and diagonal modes asymmetrically. Residual odd contributions in $\vec{k}$ no longer cancel.

Numerical Cherenkov resonances. The denominator vanishes exactly on the manifold

$$r^2 S^2(\vec{k}) = \sin^2\left(\frac{\vec{k} \cdot \vec{v} \Delta t}{2}\right) \quad (11)$$

which describes Cherenkov-radiating modes: lattice plane waves whose phase velocity equals $\vec{v}$ along the direction of motion. For any $\vec{v} \neq 0$ this manifold is nonempty in the BZ. The integrand has a $1/\sin$ -type pole on this manifold; the integral is finite only in a principal-value or causal-prescription sense (in steady state the damping ring is what regularizes it). The Cherenkov contribution gives a net force on the particle whose sign depends on which side of each pole the integration sweeps.

3.5 Predicted sign structure of $\vec{F}^{(\text{self})}(\vec{v})$

As $|\vec{v}|$ varies, the Cherenkov manifold (equation 11) reshapes inside the BZ. When $|\vec{v}|$ crosses a value at which a NEW family of modes becomes Cherenkov-resonant (i.e., the manifold acquires a new connected component or crosses a high-symmetry edge of the BZ), the net residue changes sign. The model therefore predicts DISCRETE v -thresholds at which $\vec{F}^{(\text{self})}$ sign-flips, separated by intervals of definite sign.

This matches the Stage 37e/f findings: signs change at $v \approx 0.045$ and $v \approx 0.075$ on our default grid ($c = 1$, $\Delta x = 1$, $r = 1/\sqrt{2}$), and the locations are kernel-INDEPENDENT (the pole condition does not involve $\hat{S}$).

3.6 Kernel as a low-pass filter on the integrand

The factor $|\hat{S}(\vec{k})|^2$ in equation 10 weights the integrand but does not move its poles. Wider kernels (BUMP $R = 8$, TSC + N=16 smoothing) have $\hat{S}$ that decays faster at high $|\vec{k}|$; they suppress contributions from BZ-edge modes more strongly. The net effect on the integral is to REDUCE MAGNITUDE proportionally to the kernel's high- $\vec{k}$ suppression while leaving the topology of the sign pattern unchanged. This explains the Stage 37e observation that all kernels give the same sign at every v in the scan, with magnitudes varying by factors of ~ 5 –10.

3.7 Lattice point-group symmetry of $\vec{F}^{(\text{self})}(v_x, v_y)$

The square grid has point group D_4 (4-fold rotation + reflections). $S^2(\vec{k})$ and the BZ are invariant under all D_4 operations, so equation 10 inherits the same symmetry under simultaneous

transformation of $\vec{v}$ and $\vec{k}$. This gives concrete constraints on $\vec{F}^{(\text{self})}(v_x, v_y) = (F_x, F_y)$:

$$F_x(-v_x, v_y) = -F_x(v_x, v_y) \quad (12)$$

$$F_x(v_x, -v_y) = +F_x(v_x, v_y) \quad (13)$$

$$F_x(v_y, v_x) = F_y(v_x, v_y) \quad (14)$$

$$F_x(-v_y, -v_x) = -F_y(v_x, v_y) \quad (15)$$

Equivalently $\vec{F}^{(\text{self})}$ is the gradient (up to sign) of a D_4 -invariant scalar function $U(\vec{v})$. An $\epsilon(\vec{v})$ fitted to cancel $\vec{F}^{(\text{self})}$ will inherit the same symmetries.

3.8 Continuum reference and the limit $\Delta t, \Delta x \rightarrow 0$

In the limit $\Delta t, \Delta x \rightarrow 0$ with r held fixed, the boxed integrand reduces to the continuum

$$\frac{i\vec{k}|\hat{S}(\vec{k})|^2}{c^2k^2 - (\vec{k} \cdot \vec{v})^2} \quad (16)$$

integrated over all of $\mathbb{R}^2$, which vanishes by rotation symmetry (the D_4 symmetry promotes to full $SO(2)$). All sign-flips, Cherenkov poles, and kernel-magnitude effects are discretization artifacts that vanish under joint $(\Delta t, \Delta x)$ refinement at fixed r .

4 Implementation strategy

4.1 Phase 1: per-particle field plumbing

Build the architecture with $\vec{\epsilon} = \vec{0}$ default. This gives perfect self-force cancellation IF the per-particle field's deposit + FDTD evolution faithfully reproduces what the collective field contains from that particle. Specifically:

1. Add a per-particle field-set structure (6 fields $\times$ 3 time slices + 3 source arrays per opted-in particle).
2. Modify the deposit phase: when particle p has `self_field_enabled`, deposit into BOTH the collective sources AND p 's private sources. Same kernel, same Esirkepov path.
3. Modify the FDTD step: loop the leapfrog over (collective + each enabled self-field set). Apply identical damping/PML to each.
4. Modify the force gather: if p has self-field, compute $\vec{F}_p$ per equation 3.
5. Public API for opt-in/opt-out and $\vec{\epsilon}$ setting.

4.2 Phase 2: validate that $\vec{\epsilon} = \vec{0}$ kills the spurious force

Repeat Stage 37e (single charge at constant v) with self-field enabled and $\vec{\epsilon} = 0$. Expected: $\vec{F}^{(\text{total})}(\vec{v}) = 0$ at every v in the scan (to within float32 noise). If a residual remains, characterize it and use it to define an empirical $\vec{\epsilon}(\vec{v})$.

4.3 Phase 3 (optional): tabulate the BZ integral

Numerically evaluate equation 10 on a fine $\vec{k}$ -mesh for each production kernel ($|\hat{S}(\vec{k})|^2$ for TSC, TSC+smoothing, BUMP $R = 6/8$). Compare to the measured residual from Phase 2. Agreement would validate the dispersion-based picture; disagreement would point to a missing piece in the model (likely related to the damping ring's modification of the steady-state Cherenkov regularization).

4.4 Phase 4: empirical $\bar{\epsilon}(\vec{v})$ table or closed form

If residuals persist, fit $\bar{\epsilon}(\vec{v})$ to whatever functional form captures the lattice D_4 symmetry. Candidate forms include rational functions in $\sin^2(v_x \Delta t / 2\Delta x)$, $\sin^2(v_y \Delta t / 2\Delta x)$ (the natural lattice analogs of v^2).

4.5 Backward compatibility

When all particles have `self_field_enabled = 0` (the default), the gather collapses to the existing single-collective-field behavior. Existing tests, scenarios, and WASM/web frontend code path remain unchanged.

5 Empirical results

Phase 1 of the implementation landed in commits 8fab91b (architecture) and e8e0ab0 (binary validation). The following measurements were taken with the v39 production setup: BUMP $R = 8$, equal-distance methodology, mass = 1, $G_{\text{eff}} \cdot m^2 = k_e \cdot Q^2 = 10^{-4}$ to give identical force magnitudes for the EM and GR cases.

5.1 Stage 51: single-particle uniform motion

Single $+Q$ charge in vacuum at constant v in the asymmetric 128×1024 domain. In continuum Maxwell the spurious self-force is identically zero; on the discrete grid it is the $F_{\text{spurious}}(v)$ of Stage 37e. Three configurations:

- (A) Self-field DISABLED – baseline.
- (B) Self-field ENABLED, $\epsilon = 0$ – full subtraction.
- (C) Self-field ENABLED, $\epsilon = -1$ – no subtraction.

v/c	(A) off	(B) $\epsilon = 0$	(C) $\epsilon = -1$
0.005	-8.28×10^{-7}	$+0.00 \times 10^0$	-8.28×10^{-7}
0.010	-7.74×10^{-7}	$+0.00 \times 10^0$	-7.74×10^{-7}
0.020	-9.53×10^{-7}	$+0.00 \times 10^0$	-9.53×10^{-7}
0.040	-1.17×10^{-6}	$+0.00 \times 10^0$	-1.17×10^{-6}
0.050	$+2.78 \times 10^{-7}$	$+0.00 \times 10^0$	$+2.78 \times 10^{-7}$
0.075	-2.50×10^{-6}	$+0.00 \times 10^0$	-2.50×10^{-6}
0.100	-4.81×10^{-6}	$+0.00 \times 10^0$	-4.81×10^{-6}

The cancellation in column (B) is BIT-EXACT, not just float-noise floor. This is significant: it means the per-particle FDTD evolution reproduces the collective’s contribution from that particle to floating-point precision, and the analytic $\epsilon(\vec{v})$ from the BZ integral (section 3) is NOT needed for basic cancellation. $\epsilon(\vec{v})$ becomes relevant only if one wants to leave some self-force in deliberately (radiation reaction hook or related applications).

Column (C) confirms $\epsilon = -1$ recovers the original behavior bit-exactly: $(1 + \epsilon) = 0$ means no subtraction, so the gather reduces to the collective-only path.

5.2 Stage 52: binary orbit (EM vs GR side-by-side)

The Stage 50 binary setup – both particles active sources, opposite charges (EM) or equal masses (GR), $v_{\text{orb}} = 0.01$, $r = 20$ – run for 4 orbits with self-field on/off comparison:

config	d range	$ P_x _{\max}$	$\Delta E/ E_0 $	$\Delta L_z/ L_0 $
EM self=off	[39.60, 45.28]	2.0×10^{-4}	$\pm 0.5\%$	$\pm 1.6\%$
EM self=on	[39.41, 42.74]	2.9×10^{-6}	$\pm 0.2\%$	$\pm 0.75\%$
GR self=off	[39.96, 40.07]	1.1×10^{-4}	$\pm 0.12\%$	$\pm 0.45\%$
GR self=on	[39.37, 42.71]	8.6×10^{-6}	$\pm 0.21\%$	$\pm 0.75\%$

Wins from self-field subtraction.

- Linear momentum conservation improves $50\text{--}100\times$ in both EM and GR (down to $\sim 10^{-6}$, near the float32 noise floor for a 71K-step run).
- EM rosette amplitude reduces from 13% to 7% – about half of the rosette was from the spurious self-force.
- With self-field on, EM and GR converge to ESSENTIALLY IDENTICAL numerical behavior: d ranges match to 0.05 cells, ΔE and ΔL_z match to 0.01%. This is exactly what self-field subtraction was designed to give: any remaining differences are physical (charge sign vs mass sign), not numerical.

The 7% residual rosette. GR self=off looked deceptively clean ($d \in [39.96, 40.07]$, 0.3% range). Comparing to self=on shows that was a LUCKY CANCELLATION between the spurious self-force and some other IC-driven perturbation. With self-force removed, both EM and GR show the same $\sim 7\%$ rosette amplitude – this is now the cleanest measurement of the non-self-force IC amplification under Bertrand’s theorem on the 2D-log potential. Sources of the residual rosette include (a) the initial-condition $\vec{x}_p(0)$ sitting at the COM grid intersection (which has sub-cell offsets), (b) the bipolar ρ_q float-roundoff effect in EM (no GR analog, but EM and GR converge to the same rosette anyway, which constrains this), and (c) physical residual from FDTD’s finite- c wake formation during the first few orbital periods. Not blocking for production but worth understanding in a future stage if precision matters.

The bit-exact cancellation principle. Combining the Stage 51 bit-exact zero with the Stage 52 $50\text{--}100\times$ Px improvement, the v39 architecture’s claim “the per-particle FDTD correctly captures the collective’s per-particle contribution” is validated empirically. Analytical $\epsilon(\vec{v})$ work (section 3) becomes a tool for OTHER purposes – analytic understanding, future radiation-reaction module, or phenomenological extensions – not a requirement for basic self-force cancellation.

6 Closing the rosette: IC calibration to the discrete force law

The $\sim 7\%$ residual rosette in Stage 52 is investigated in Stages 53 (c-scan) and 53b (v_{orb} -factor scan).

6.1 c-scan rules out delayed potential

Hypothesis (to be tested): the rosette arises from each particle responding to the OTHER particle’s retarded position. At $t = 0$ the field is zero and the mutual force takes $\tau = d/c$ to build up; during that time each particle drifts radially outward (no centripetal force), giving an off-circular IC that Bertrand’s theorem amplifies into a rosette.

Test: vary c_{eff} at fixed v_{orb} . Since $v_{\text{orb}} = Q\sqrt{k_e/m}$ does not depend on c , doubling c halves τ while leaving the orbital period unchanged. Smaller τ/T should correspond to smaller field-buildup perturbation and hence smaller rosette.

config	c_{eff}	τ/T	d range
EM	1	0.00318	7.06%
EM	2	0.00159	7.08%
EM	4	0.00080	7.33%
GR	1	0.00318	7.05%
GR	2	0.00159	7.08%
GR	4	0.00080	7.32%

The rosette amplitude is C-INDEPENDENT to better than 0.4% across a $4\times$ change in c . Hypothesis falsified.

6.2 v_{orb} scan locates the discrete-circular velocity

The scenario's v_{orb} comes from the CONTINUUM 2D log Coulomb force law:

$$m \frac{v_{\text{orb}}^2}{r} = \frac{k_e Q^2}{r} \implies v_{\text{orb}} = Q \sqrt{k_e/m}. \quad (17)$$

The DISCRETE force on each particle from the OTHER's kernel-smeared deposit at separation d differs from $k_e Q^2/d$ by a kernel- and separation-dependent correction. For BUMP $R = 8$ at $d = 40$ (kernel half-width ~ 8 , separation/kernel ratio ~ 2.4), the correction is of order few percent. Using the continuum velocity as IC then puts the orbit on a slightly non-circular state that Bertrand amplifies into the visible rosette.

Test: scan v_{factor} (multiplier on continuum v_{orb}) and measure the rosette range. Self-field is ON throughout (so the spurious self-force is removed and we're seeing only the IC mismatch).

v_{factor}	d range	d_{min}	d_{max}
0.95	5.15%	37.95	40.01
0.97	2.12%	39.43	40.28
0.98	1.31%	40.00	40.52
0.99	3.42%	40.00	41.37
1.000	7.06%	39.71	42.53
1.005	7.73%	39.85	42.94
1.010	8.01%	40.00	43.20
1.020	9.83%	40.00	43.93
1.030	14.08%	39.65	45.28
1.050	17.80%	40.00	47.12

The minimum rosette amplitude is at $v_{\text{factor}} \approx 0.98$ (range 1.31%), not at 1.0. The discrete-circular velocity is $\sim 2\%$ lower than the continuum prediction.

Interpretation: $v_{\text{factor}}^2 = F_{\text{discrete}}/F_{\text{continuum}} \approx 0.96$ gives $v_{\text{factor}} \approx 0.98$, i.e., the actual mutual force at this separation is about 4% smaller than the point-source continuum value. This is consistent with the kernel-induced smoothing of the deposit reducing the effective coupling at finite d/R_{kernel} .

6.3 Synthesis: rosette source breakdown

Combining the chain of falsifications and tests across Stages 47–53b:

Source	Contribution to original 13% drift	Status
Spurious self-force	$\sim 6\%$	Removed by v39 self-fie
IC velocity mismatch (continuum vs discrete)	$\sim 7\%$	Removed by setting v_{fa}
Delayed potential	$\sim 0\%$	Falsified by c-sc
Bipolar ρ_q float roundoff (EM-only)	negligible	Constrained by EM/GR

6.4 Recipe for effectively circular orbits in v39

1. Use the v38 production setup (BUMP $R = 8$ or $R = 6$, LB gather, PML damping).
2. Enable self-field on every particle that participates as an active source (Stage 51 shows bit-exact spurious-force cancellation at $\epsilon = 0$).
3. Calibrate the IC orbital velocity to the discrete-circular value, which differs from the continuum prediction by a few percent depending on kernel choice and d/R_{kernel} . For BUMP $R = 8$ at $d = 40$ the calibration factor is ~ 0.98 ; other configurations should be tabulated similarly.

With all three in place, the binary orbit ranges over 1.3% in d over 2 orbits at the discrete-circular IC – essentially circular within float32 precision.

This closes the chain of self-force / rosette investigations that have run through Stages 37–53b. Remaining v/c thresholds in the single-particle uniform-motion case (Stage 37e) are the dispersion-driven Cherenkov features documented in section 3. They cannot be eliminated by kernel choice and are intrinsic to FDTD on a fixed grid; the self-field subtraction route is the principled mitigation.